Word Games

J Chiu

September 22, 2021

Natural Language Interaction

- Interaction (through language) is important
  - Cannot fully automate every task, i.e. task-oriented or information seeking dialogues require human input
  - Must handle diverse non-expert human input, although input may map to a low-dimensional manifold
  - High levels of ambiguity must be resolved via interaction
- Interaction (through language) is hard
  - Human input is expensive, so supervision is limited
  - In order to make certain problem aspects tractable, must make sacrifices in other areas (toy domain $=$ out of distribution for pretrained models)
- What are the main challenges in interaction, and what are the tradeoffs of different approaches?

Types of Dialogue Games

- Task-oriented: Wizard of Oz (WoZ)
    - Tseng et al. (2019): Wizard obtains task from human then executes it.
- Deliberation / reference / signal
    - Udagawa and Aizawa (2019): Visual reference game with latent translated views. Each player gets a different petri dish view of the same underlying game board, and players must select the same object on the board.
- Information seeking / inquiry
    - Yu et al. (2019): WoZ-style answer providing where asker does not know exact question. Latent true question (to all), WoZ must answer
- Persuasion / negotiation
    - Lewis et al. (2017): Negotiation over an observed set of item with latent utilities for each agent.

Types of Dialogue Games

In all cases, the game can be (indirectly) solved by resolving a latent variable

- ▶ When is this tractable, and why do no new methods do this?
- ▶ New (ie basically all) methods rely on supervision
- ▶ If they do not, it is because the game has a trivial solution

Types of Dialogue Games: Latent goals and strategy

What are the latent variables in each type?

- ▶ Task-oriented
  - ▶ Latent task slots
- ▶ Deliberation / reference / signal
  - ▶ Varies per game
- ▶ Information seeking / inquiry
  - ▶ Infer true question, find answer
- ▶ Persuasion / negotiation
  - ▶ Infer utilities, exploit

Types of Dialogue Games: Latent goals and strategy

▶ Tasks must be interesting enough so that latent quantities cannot be inferred with a single utterance, reducing them to single turn games
  ▶ High degree of ambiguity / distractors or large number of slots to fill (combinatorial)
▶ Break down latent quantities and use heuristics to make assumptions on structure
  ▶ For example, choosing an ordering of WoZ slots: When choosing a restaurant, first figure out time, then cuisine, and finally price
  ▶ Will likely remain task-specific
▶ What other parts can we learn?

Belief State Tracking

- ▶ The incremental inference procedure is known as belief state tracking (BST)
- ▶ Local semantics are aggregated into belief state, which informs high-level strategic decisions
- ▶ Seems difficult to learn language, high level strategy, belief state updating, and low level parsing at the same time
- ▶ Ablate how structure influences each of these

Language Games

- Games offer a testbed for the development of methods
  - Allow designers to control difficulty and simplicity
- Allowing interaction through language increases the population of players
-

3 Challenges in Language Games

1. Game design
   - ▶ Simplification of or step towards a realistic task
   - ▶ Must be fun or useful
2. Meaning representations
   - ▶
3. Search
   - ▶ A
4. Modeling other agents
   - ▶ A

Axes of Tractability

- Horizon length
  - B
- Action space
  - B
- A
  - B

Meaning reps

- In full generality, this problem is often encountered in hierarchical RL
    - Less bleak in the language gamesetting
    - Games are often very simple and can be constrained to small horizons, for example He He engineered a parser and policy that basically solves the negotiation task
- Many text-specific meaning representations (MR) to choose from
    - Many are too complex
    - Can we leverage existing MRs to learn a minimal task-specific representation that balances utility and expressivity?

Learning to plan

▶

Contributions

▶ Under a unified Bayesian game perspective of dialog games, present formulations for different classes of dialogs: signaling, negotiation?

▶ Provide pipelined variational Bayesian framework for learning to play dialog games from offline data, with and without granular annotations

▶ Good results

Tasks

- Task-oriented dialogue: Multi-WoZ
- Negotation: Deal-or-no-deal
- Reference: OneCommon
- Information-Seeking: Birds

Citations I

Lewis, M., Yarats, D., Dauphin, Y. N., Parikh, D., and Batra, D. (2017). Deal or no deal? end-to-end learning for negotiation dialogues. *CoRR*, abs/1706.05125.

Tseng, B., Rei, M., Budzianowski, P., Turner, R. E., Byrne, B., and Korhonen, A. (2019). Semi-supervised bootstrapping of dialogue state trackers for task oriented modelling. *CoRR*, abs/1911.11672.

Udagawa, T. and Aizawa, A. (2019). A natural language corpus of common grounding under continuous and partially-observable context. *CoRR*, abs/1907.03399.

Yu, L., Chen, H., Wang, S. I., Artzi, Y., and Lei, T. (2019). Interactive classification by asking informative questions. *CoRR*, abs/1911.03598.